

Running perceptron on the airline satisfaction dataset and experimenting with kernel methods using an implementation of kernelized perceptron from scratch and kernel approximation.
@arashsm79

Contents:

- [Pre-processing](#)
- [1- Perceptron](#)
- [2- Non-linear Perceptron](#)
 - [Kernel Trick](#)
 - [Kernel Approximation](#)

```
In [1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Pre-processing

Before creating a model from the given data, I first have to clean it and get it ready for further processing in the perceptron. I will explain each step one by one.

```
In [2]: # Import the raw data as a pandas data frame
raw_train_data = pd.read_csv('train.csv')
```

```
In [3]: raw_train_data.head()
```

```
Out[3]:
```

	Unnamed: 0	id	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	...	Inflight entertainment	On-board service	Leg room service	Baggage handling
0	0	70172	Male	Loyal Customer	13	Personal Travel	Eco Plus	460	3	4	...	5	4	3	
1	1	5047	Male	disloyal Customer	25	Business travel	Business	235	3	2	...	1	1	5	
2	2	110028	Female	Loyal Customer	26	Business travel	Business	1142	2	2	...	5	4	3	
3	3	24026	Female	Loyal Customer	25	Business travel	Business	562	2	5	...	2	2	5	
4	4	119299	Male	Loyal Customer	61	Business travel	Business	214	3	3	...	3	3	4	

5 rows × 25 columns

```
In [4]: raw_train_data.shape
```

```
Out[4]: (103904, 25)
```

```
In [5]: raw_train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            103904 non-null int64
1   id                                    103904 non-null int64
2   Gender                                103904 non-null object
3   Customer Type                         103904 non-null object
4   Age                                    103904 non-null int64
5   Type of Travel                        103904 non-null object
6   Class                                103904 non-null object
7   Flight Distance                       103904 non-null int64
8   Inflight wifi service                 103904 non-null int64
9   Departure/Arrival time convenient    103904 non-null int64
10  Ease of Online booking                103904 non-null int64
11  Gate location                         103904 non-null int64
12  Food and drink                        103904 non-null int64
13  Online boarding                       103904 non-null int64
14  Seat comfort                          103904 non-null int64
15  Inflight entertainment                103904 non-null int64
16  On-board service                      103904 non-null int64
17  Leg room service                      103904 non-null int64
18  Baggage handling                      103904 non-null int64
19  Checkin service                       103904 non-null int64
20  Inflight service                      103904 non-null int64
21  Cleanliness                           103904 non-null int64
22  Departure Delay in Minutes            103904 non-null int64
23  Arrival Delay in Minutes              103594 non-null float64
24  satisfaction                           103904 non-null object
dtypes: float64(1), int64(19), object(5)
memory usage: 19.8+ MB
```

First we need to separate the useless columns from the actual data. The first and the second columns seem to have no information of consequence, thus, I will remove them. I also need to separate the labels from the actual data.

```
In [6]: train_data = raw_train_data[raw_train_data.columns[2:-1]].copy()
        train_data_label = raw_train_data['satisfaction']
```

From the above information, we can see that 'Arrival Delay in Minutes' has some null values; I will have to replace them with a proper value. (here I chose the mean value)

```
In [7]: train_data['Arrival Delay in Minutes'].replace(np.nan, train_data['Arrival Delay in Minutes'].mean(), inplace=True)
```

Normalize numerical data into a range from 0 to 1.

Here I have used MinMaxScale which transforms features by scaling each feature to a given range. This estimator scales and translates each feature individually such that it is in the given range on the training set, e.g. between zero and one. The transformation is given by:

$$X_std = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))$$

$$X_scaled = X_std * (max - min) + min$$

Advantages of normalizing the data:

- speeds up learning and leads to faster convergence
- changes the values of numeric columns in the dataset to a common scale
- helps linear seerators such as SVMs and perceptrons

```
In [8]: from sklearn.preprocessing import MinMaxScaler

num_attributes = train_data.select_dtypes(include=['float64', 'int64']).columns.tolist()
minmax_sc = MinMaxScaler()
minmax_sc.fit(train_data[num_attributes])
train_data[num_attributes] = minmax_sc.transform(train_data[num_attributes])
```

It is also possible to use other scalers such as StandardScaler.

```
In [9]: # from sklearn.preprocessing import StandardScaler
        # sc = StandardScaler()
        # train_data[num_attributes] = sc.fit_transform(train_data[num_attributes])
```

I need to some how encode categorical data for further processing. Here I have chosen OneHotEncoder, because it seems to be the standard for linear models such as SVM and perceptron.

The features are encoded using a one-hot encoding scheme. This creates a binary column for each category and returns a sparse matrix or dense array

By default, the encoder derives the categories based on the unique values in each feature.

This encoding is needed for feeding categorical data to many scikit-learn estimators, notably linear models and SVMs with the standard kernels.

```
In [10]: cat_attributes = train_data.select_dtypes(include=['object']).columns.tolist()
train_data = pd.get_dummies(train_data, columns=cat_attributes)
```

```
In [11]: train_data.head()
```

```
Out[11]:
```

	Age	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	Ease of Online booking	Gate location	Food and drink	Online boarding	Seat comfort	Inflight entertainment	...	Arrival Delay in Minutes	Gender_Female	Gender_Male
0	0.076923	0.086632	0.6	0.8	0.6	0.2	1.0	0.6	1.0	1.0	...	0.011364	0	1
1	0.230769	0.041195	0.6	0.4	0.6	0.6	0.2	0.6	0.2	0.2	...	0.003788	0	1
2	0.243590	0.224354	0.4	0.4	0.4	0.4	1.0	1.0	1.0	1.0	...	0.000000	1	0
3	0.230769	0.107229	0.4	1.0	1.0	1.0	0.4	0.4	0.4	0.4	...	0.005682	1	0
4	0.692308	0.036955	0.6	0.6	0.6	0.6	0.8	1.0	1.0	0.6	...	0.000000	0	1

5 rows × 27 columns

Now we repeat the above steps for test data.

```
In [12]: # load
raw_test_data = pd.read_csv('test.csv')
test_data = raw_test_data[raw_test_data.columns[2:-1]].copy()
test_data_label = raw_test_data['satisfaction']

# remove nan
test_data['Arrival Delay in Minutes'].replace(np.nan, test_data['Arrival Delay in Minutes'].mean(), inplace=True)

# normalize
test_num_attributes = test_data.select_dtypes(include=['float64', 'int64']).columns.tolist()
minmax_sc = MinMaxScaler()
minmax_sc.fit(test_data[test_num_attributes])
test_data[test_num_attributes] = minmax_sc.transform(test_data[test_num_attributes])

# encode categorical data
test_cat_attributes = test_data.select_dtypes(include=['object']).columns.tolist()
test_data = pd.get_dummies(test_data, columns=test_cat_attributes)
```

```
In [13]: test_data.head()
```

```
Out[13]:
```

	Age	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	Ease of Online booking	Gate location	Food and drink	Online boarding	Seat comfort	Inflight entertainment	...	Arrival Delay in Minutes	Gender_Female	Gender_Male
0	0.576923	0.026050	1.0	0.8	0.6	0.75	0.6	0.8	0.50	1.0	...	0.039462	1	0
1	0.371795	0.571890	0.2	0.2	0.6	0.00	1.0	0.8	1.00	0.8	...	0.000000	1	0
2	0.166667	0.032512	0.4	0.0	0.4	0.75	0.4	0.4	0.25	0.4	...	0.000000	0	1
3	0.474359	0.675687	0.0	0.0	0.0	0.25	0.6	0.8	0.75	0.2	...	0.005381	0	1
4	0.538462	0.232431	0.4	0.6	0.8	0.50	0.8	0.2	0.25	0.4	...	0.017937	1	0

5 rows × 27 columns

1-Perceptron

```
In [14]: from sklearn.linear_model import Perceptron
```

```
In [15]: p = Perceptron()
p.fit(train_data, train_data_label)
```

```
Out[15]: Perceptron()
```

```
In [16]: from sklearn.metrics import accuracy_score

# predict on train
pred_train = p.predict(train_data)
train_score = accuracy_score(train_data_label, pred_train)
print("Train data accuracy: ", "{:.2f}%".format(train_score*100))

# predict on test
pred_test = p.predict(test_data)
test_score = accuracy_score(test_data_label, pred_test)
print("Test data accuracy: ", "{:.2f}%".format(test_score*100))
```

Train data accuracy: 82.11%
Test data accuracy: 82.78%

2-Non-linear Perceptron

Kernel Trick

There are no standard Perceptron implementations that use the kernel trick. Thus, I have to implement a binary classifier perceptron from scratch.

I have used the algorithm described in the following resources:

- Shawe-Taylor, John; Cristianini, Nello (2004). Kernel Methods for Pattern Analysis. Cambridge University Press. pp. 241–242.
- https://en.wikipedia.org/wiki/Kernel_perceptron
- <https://webpages.charlotte.edu/rbunescu/courses/ou/ml4900/lecture06.pdf>
- https://alex.smola.org/teaching/pune2007/pune_3.pdf

The gist of it is that, we form a dual problem and create a kernelized perceptron algorithm that uses the dot products of the training samples. Then we use the `kernel trick` to compute the dot product in the higher dimension using low dimension data.

Since this is a binary classification (satisfied or dissatisfied/neutral, we don't need to worry about handling multiple classes using methods such as one-vs-one or one-vs-all).

This code is written in python (in contrast to library functions of sklearn that are implemented in C) and thus, using the whole train data takes an infeasible amount of time.

```
In [49]: reduced_train_data = train_data.head(1000).copy()
reduced_train_data_label = train_data_label.head(1000).copy()
reduced_train_data_label = pd.get_dummies(reduced_train_data_label, columns='satisfaction', drop_first=True)

reduced_test_data = test_data.head(1000).copy()
reduced_test_data_label = test_data_label.head(1000).copy()
reduced_test_data_label = pd.get_dummies(reduced_test_data_label, columns='satisfaction', drop_first=True)
```

Polynomial Kernels in \mathbb{R}^n

Idea

- We want to extend $k(x, x') = \langle x, x' \rangle^2$ to

$$k(x, x') = (\langle x, x' \rangle + c)^d \text{ where } c > 0 \text{ and } d \in \mathbb{N}.$$

- Prove that such a kernel corresponds to a dot product.

Proof strategy

Simple and straightforward: compute the explicit sum given by the kernel, i.e.

$$k(x, x') = (\langle x, x' \rangle + c)^d = \sum_{i=0}^d \binom{d}{i} (\langle x, x' \rangle)^i c^{d-i}$$

Individual terms $(\langle x, x' \rangle)^i$ are dot products for some $\Phi_i(x)$.

```
In [43]: def polynomial_kernel(X1, X2, c=1, d=2):
         return (np.dot(X1, X2) + c) ** d
```

```
In [44]: alpha = []
support_vectors = []
support_vectors_y = []
epsilon = 1e-10
T = 2
kernel = polynomial_kernel
```

```
In [45]: # generate the gram matrix
def gram_matrix(X1, X2 = None):
    if X2 is None:
        X2 = X1
```

```

n, _ = X1.shape
m, _ = X2.shape
K = np.zeros((n, m))
for i in range(n):
    for j in range(m):
        K[i, j] = kernel(X1[i], X2[j])
return K

```

1. **define** $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} = \sum_n \alpha_n \mathbf{x}_n^T \mathbf{x} = \sum_n \alpha_n K(\mathbf{x}_n, \mathbf{x})$
2. **initialize** dual parameters $\alpha_n = 0$
3. **for** $n = 1 \dots N$
4. $h_n = \text{sgn } f(\mathbf{x}_n)$
5. **if** $h_n \neq t_n$ **then**
6. $\alpha_n = \alpha_n + t_n$

During testing: $h(\mathbf{x}) = \text{sgn } f(\mathbf{x})$

```

In [46]: def fit(train_data_x, train_data_y):
    global alpha
    global support_vectors
    global support_vectors_y
    n, dim = train_data_x.shape
    alpha = np.zeros(n, dtype=np.float64)

    K = gram_matrix(train_data_x)

    for t in range(T):
        for i in range(n):
            if np.sign(np.sum(alpha * train_data_y * K[:, i])) != train_data_y[i]:
                alpha[i] += 1

    support_vectors_mask = alpha > epsilon
    alpha = alpha[support_vectors_mask]
    support_vectors = train_data_x[support_vectors_mask]
    support_vectors_y = train_data_y[support_vectors_mask]

```

```

In [47]: def predict(X):
    y_predict = np.zeros(len(X))
    global alpha
    global support_vectors
    global support_vectors_y
    for i in range(len(X)):
        sum = 0
        for j in range(len(alpha)):
            curr_alpha = alpha[j]
            support_vector_y = support_vectors_y[j]
            support_vector = support_vectors[j]
            sum += curr_alpha * support_vector_y * \
                kernel(X[i], support_vector)
        y_predict[i] = sum
    return np.sign(y_predict)

```

```

In [50]: fit(reduced_train_data.to_numpy(), reduced_train_data_label.to_numpy())
pred_test = predict(reduced_test_data.to_numpy())
test_score = accuracy_score(reduced_test_data_label, pred_test)
print("Test data accuracy: ", "{:.2f}%".format(test_score*100))

```

Test data accuracy: 45.50%

the above score is only for 1000 instances of data! for the whole set, it will be much higher.

Kernel Approximation

Here is what I learned from reading the sklearn documentations about this approach:

Kernel approximation performs non-linear transformations of the input, which can serve as a basis for linear classification or other algorithms.

The advantages of using Kernel approximation compared to the kernel trick:

- can be better suited for online learning
- can significantly reduce the cost of learning with very large datasets
- Standard kernelized SVMs do not scale well to large datasets, but using an approximate kernel map it is possible to use much more efficient linear SVMs.

2 methods are described in the sklearn documentations:

- Method for Kernel Approximation
- RadialNystroeml Basis Function Kernel

RBF

```
In [82]: from sklearn import pipeline
from sklearn.kernel_approximation import RBFSampler

# creating a pipeline
rbf_map = RBFSampler(gamma=0.04, random_state=42)
rbf_approx_perceptron = pipeline.Pipeline([("rbf", rbf_map), ("Perceptron", Perceptron())])

# fitting on train data
rbf_approx_perceptron.fit(train_data, train_data_label)

# predict test and compute score
pred_test = rbf_approx_perceptron.predict(test_data)
test_score = accuracy_score(test_data_label, pred_test)
print("Test data accuracy: ", "{:.2f}%".format(test_score*100))
```

Test data accuracy: 90.56%

Nystroem

```
In [72]: from sklearn.kernel_approximation import Nystroem

nystroem_map = Nystroem(gamma=0.1, random_state=42)
nystroem_approx_perceptron = pipeline.Pipeline([("nystroem", nystroem_map), ("Perceptron", Perceptron())])

# fitting on train data
nystroem_approx_perceptron.fit(train_data, train_data_label)

# predict test and compute score
pred_test = nystroem_approx_perceptron.predict(test_data)
test_score = accuracy_score(test_data_label, pred_test)
print("Test data accuracy: ", "{:.2f}%".format(test_score*100))
```

Test data accuracy: 91.37%